

Importing Libraries

```
In [70]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [71]: df=pd.read_csv("madrid_2013.csv").fillna(1)
df
```

Out[71]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2013-11-01 01:00:00	1.0	0.6	1.0	1.0	135.0	74.0	1.0	1.0	1.0	7.0	1.0	1.0	28
1	2013-11-01 01:00:00	1.5	0.5	1.3	1.0	71.0	83.0	2.0	23.0	16.0	12.0	1.0	8.3	28
2	2013-11-01 01:00:00	3.9	1.0	2.8	1.0	49.0	70.0	1.0	1.0	1.0	1.0	1.0	9.0	28
3	2013-11-01 01:00:00	1.0	0.5	1.0	1.0	82.0	87.0	3.0	1.0	1.0	1.0	1.0	1.0	28
4	2013-11-01 01:00:00	1.0	1.0	1.0	1.0	242.0	111.0	2.0	1.0	1.0	12.0	1.0	1.0	28
...	
209875	2013-03-01 00:00:00	1.0	0.4	1.0	1.0	8.0	39.0	52.0	1.0	1.0	1.0	1.0	1.0	28
209876	2013-03-01 00:00:00	1.0	0.4	1.0	1.0	1.0	11.0	1.0	6.0	1.0	2.0	1.0	1.0	28
209877	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	2.0	4.0	75.0	1.0	1.0	1.0	1.0	1.0	28
209878	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	2.0	11.0	52.0	1.0	1.0	1.0	1.0	1.0	28
209879	2013-03-01 00:00:00	1.0	1.0	1.0	1.0	1.0	10.0	75.0	3.0	1.0	1.0	1.0	1.0	28

209880 rows × 14 columns



Data Cleaning and Data Preprocessing

In [72]: `df=df.dropna()`

In [73]: `df.columns`

Out[73]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
 'SO_2', 'TCH', 'TOL', 'station'],
 dtype='object')

In [74]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209880 non-null  object
1   BEN         209880 non-null  float64
2   CO          209880 non-null  float64
3   EBE         209880 non-null  float64
4   NMHC        209880 non-null  float64
5   NO          209880 non-null  float64
6   NO_2        209880 non-null  float64
7   O_3         209880 non-null  float64
8   PM10        209880 non-null  float64
9   PM25        209880 non-null  float64
10  SO_2        209880 non-null  float64
11  TCH         209880 non-null  float64
12  TOL         209880 non-null  float64
13  station     209880 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 24.0+ MB
```

```
In [75]: data=df[['CO' , 'station']]
data
```

Out[75]:


	CO	station
0	0.6	28079004
1	0.5	28079008
2	1.0	28079011
3	0.5	28079016
4	1.0	28079017
...
209875	0.4	28079056
209876	0.4	28079057
209877	1.0	28079058
209878	1.0	28079059
209879	1.0	28079060

209880 rows × 2 columns

Line chart

```
In [76]: data.plot.line(subplots=True)
```


Out[76]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [77]: data.plot.line()
```

```
Out[77]: <AxesSubplot:>
```




Bar chart

```
In [78]: b=data[0:50]
```

```
In [79]: b.plot.bar()
```


```
Out[79]: <AxesSubplot:>
```



Histogram

```
In [80]: data.plot.hist()
```


```
Out[80]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [81]: data.plot.area()
```


```
Out[81]: <AxesSubplot:>
```



Box chart

```
In [82]: data.plot.box()
```

```
Out[82]: <AxesSubplot:>
```



Pie chart

```
In [83]: b.plot.pie(y='station' )
```


```
Out[83]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [84]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[84]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [85]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 209880 entries, 0 to 209879
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209880 non-null  object
1   BEN         209880 non-null  float64
2   CO          209880 non-null  float64
3   EBE         209880 non-null  float64
4   NMHC        209880 non-null  float64
5   NO          209880 non-null  float64
6   NO_2        209880 non-null  float64
7   O_3         209880 non-null  float64
8   PM10        209880 non-null  float64
9   PM25        209880 non-null  float64
10  SO_2        209880 non-null  float64
11  TCH         209880 non-null  float64
12  TOL         209880 non-null  float64
13  station     209880 non-null  int64
```


In [86]: `df.describe()`

Out[86]:

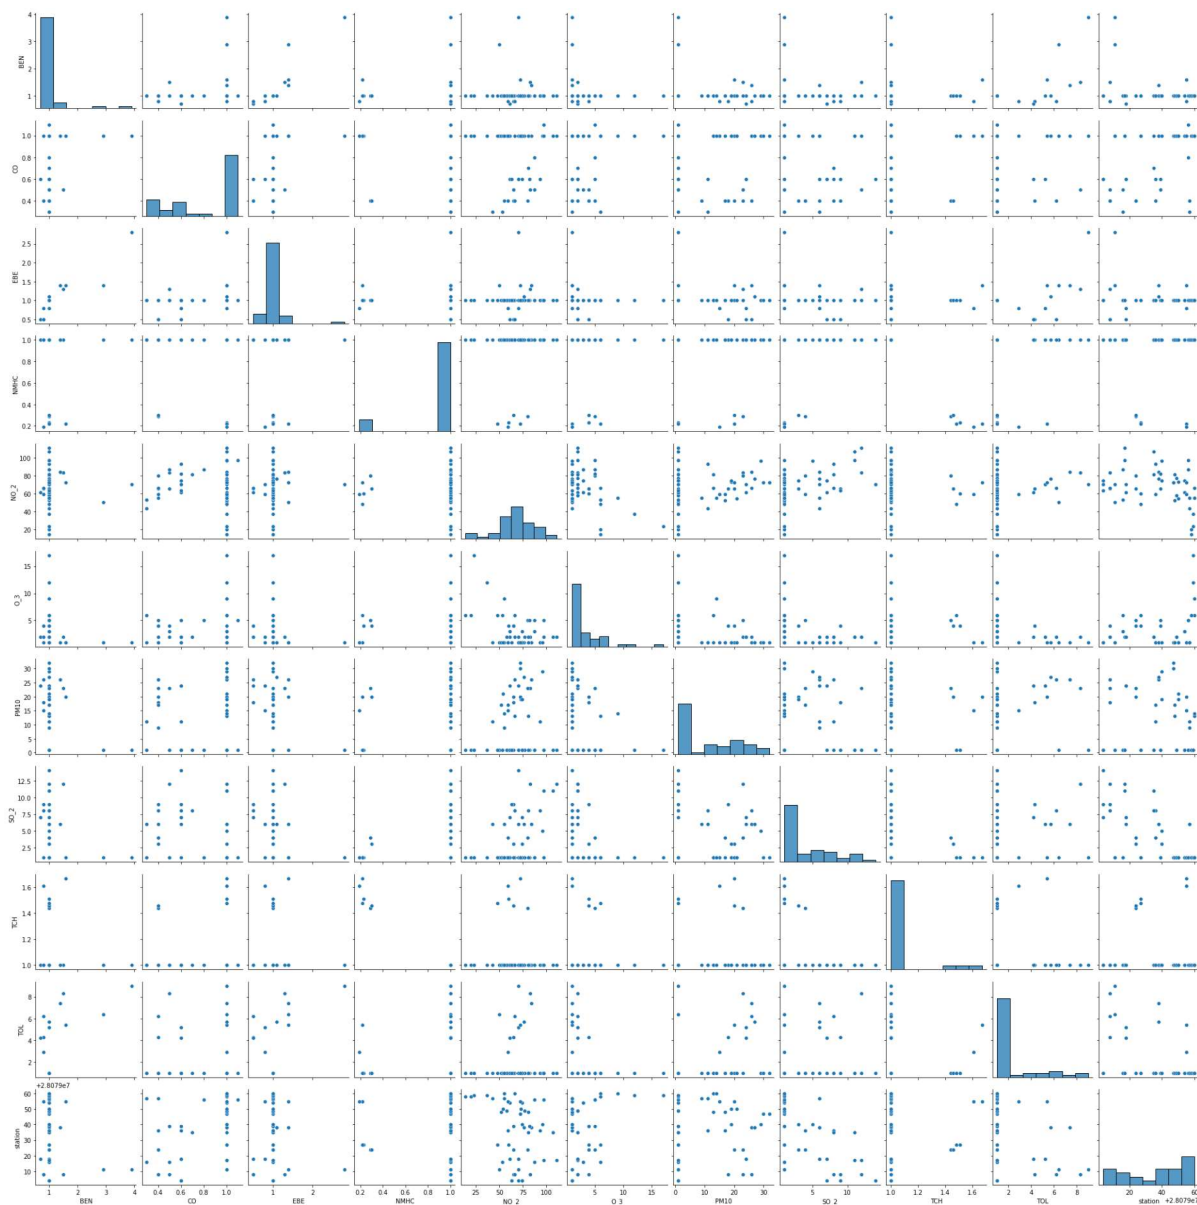
	BEN	CO	EBE	NMHC	NO	NO_
count	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000	209880.000000
mean	0.931014	0.721695	0.954744	0.900223	20.101401	34.58640
std	0.430684	0.361528	0.301074	0.267139	44.319112	27.86658
min	0.100000	0.100000	0.100000	0.040000	1.000000	1.000000
25%	1.000000	0.300000	1.000000	1.000000	2.000000	14.00000
50%	1.000000	1.000000	1.000000	1.000000	5.000000	27.00000
75%	1.000000	1.000000	1.000000	1.000000	17.000000	48.00000
max	12.100000	10.400000	11.800000	1.000000	1081.000000	388.00000

In [87]: `df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
'PM10', 'SO_2', 'TCH', 'TOL', 'station']]`

EDA AND VISUALIZATION

```
In [88]: sns.pairplot(df1[0:50])
```

```
Out[88]: <seaborn.axisgrid.PairGrid at 0x21084f226a0>
```




```
In [89]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

```
Out[89]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [90]: sns.heatmap(df1.corr())
```

```
Out[90]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [91]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
               'PM10', 'SO_2', 'TCH', 'TOL']]
          y=df['station']
```

```
In [92]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [93]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[93]: LinearRegression()

```
In [94]: lr.intercept_
```

Out[94]: 28078974.171740144


```
In [95]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[95]:

	Co-efficient
BEN	2.179739
CO	18.374955
EBE	10.083815
NMHC	18.720052
NO_2	-0.055955
O_3	0.009840
PM10	0.206652
SO_2	-0.933468
TCH	27.406878
TOL	-3.694933

```
In [96]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[96]: <matplotlib.collections.PathCollection at 0x21091bd7610>



ACCURACY

```
In [97]: lr.score(x_test, y_test)
```

Out[97]: 0.29798024782200694

```
In [98]: lr.score(x_train, y_train)
```

Out[98]: 0.3005822000263991

Ridge and Lasso

```
In [99]: from sklearn.linear_model import Ridge, Lasso
```

```
In [100]: rr=Ridge(alpha=10)
rr.fit(x_train, y_train)
```

Out[100]: Ridge(alpha=10)

Accuracy(Ridge)

```
In [101]: rr.score(x_test, y_test)
```

Out[101]: 0.29800039047740845

```
In [102]: rr.score(x_train,y_train)
```

```
Out[102]: 0.300578924013424
```

Accuracy(Lasso)

```
In [103]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[103]: Lasso(alpha=10)
```

```
In [104]: la.score(x_train,y_train)
```

```
Out[104]: 0.04519637644952712
```

ElasticNet

```
In [105]: la.score(x_test,y_test)
```

```
Out[105]: 0.04452577535469704
```

```
In [106]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[106]: ElasticNet()
```

```
In [107]: en.coef_
```

```
Out[107]: array([ 0.40223113,  2.69520561,  0.5174599 ,  0.          , -0.02040017,  
                -0.01561821,  0.16279299, -1.27581233, -0.          , -1.66911887])
```

```
In [108]: en.intercept_
```

```
Out[108]: 28079039.963054292
```

```
In [109]: prediction=en.predict(x_test)
```

```
In [110]: en.score(x_test,y_test)
```

```
Out[110]: 0.15303243515292386
```

Evaluation Metrics

```
In [111]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
13.733731326439136
262.79454649641025
16.21093909976872
```

Logistic Regression

```
In [112]: from sklearn.linear_model import LogisticRegression
```

```
In [113]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
                             'PM10', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [114]: feature_matrix.shape
```

```
Out[114]: (209880, 10)
```

```
In [115]: target_vector.shape
```

```
Out[115]: (209880,)
```

```
In [116]: from sklearn.preprocessing import StandardScaler
```

```
In [117]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [118]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[118]: LogisticRegression(max_iter=10000)
```

```
In [121]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [122]: prediction=logr.predict(observation)
print(prediction)
```

```
[28079008]
```

```
In [123]: logr.classes_
```

```
Out[123]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
                  28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
                  28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
                  28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
                dtype=int64)
```

```
In [124]: logr.score(fs,target_vector)
```

```
Out[124]: 0.6612921669525443
```

```
In [125]: logr.predict_proba(observation)[0][0]
```

```
Out[125]: 9.49253547859177e-217
```

```
In [126]: logr.predict_proba(observation)
```

```
Out[126]: array([[9.49253548e-217, 6.03969072e-001, 1.69773000e-169,
                  1.44179094e-134, 1.71060740e-074, 3.96021369e-001,
                  9.55808997e-006, 5.22717178e-089, 5.48319507e-081,
                  1.32436170e-079, 1.07294134e-076, 3.50636612e-129,
                  1.69529056e-079, 3.82520459e-158, 4.22872970e-161,
                  3.57928159e-187, 2.10845766e-164, 8.33937392e-188,
                  1.12752042e-082, 7.42692411e-129, 7.66872499e-080,
                  6.30044443e-191, 4.32093567e-191, 3.26054498e-071]])
```

Random Forest

```
In [127]: from sklearn.ensemble import RandomForestClassifier
```

```
In [128]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[128]: RandomForestClassifier()
```

```
In [129]: parameters={'max_depth':[1,2,3,4,5],
                      'min_samples_leaf':[5,10,15,20,25],
                      'n_estimators':[10,20,30,40,50]}
}
```

```
In [130]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[130]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [131]: grid_search.best_score_
```

```
Out[131]: 0.6921165836260176
```

```
In [132]: rfc_best=grid_search.best_estimator_
```



```
In [133]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'],
```

```
Out[133]: [Text(2480.0, 1993.2, 'SO_2 <= 1.5\ngini = 0.958\nsamples = 92867\nvalue =
[6087, 5963, 6200, 6122, 6172, 6108, 6139, 6071, 6153\n6021, 6094, 6162, 61
29, 6265, 6136, 6252, 6120, 6171\n6106, 6035, 5939, 6179, 5970, 6322]\nclas
s = x'),
Text(1550.0, 1630.8000000000002, 'CO <= 0.95\ngini = 0.936\nsamples = 5823
8\nvalue = [11, 49, 6200, 6122, 71, 40, 1851, 6071, 8, 3158, 7\n6162, 900,
6265, 6136, 6252, 6120, 6171, 6106, 6035\n30, 6179, 5970, 6322]\nnclass =
x'),
Text(806.0, 1268.4, 'PM10 <= 1.5\ngini = 0.773\nsamples = 14369\nvalue =
[3, 11, 0, 6017, 0, 33, 1831, 0, 6, 3156, 0, 5961\n0, 0, 0, 0, 0, 0, 574
2, 20, 0, 0, 0]\nnclass = d'),
Text(496.0, 906.0, 'CO <= 0.15\ngini = 0.668\nsamples = 11200\nvalue = [3,
0, 0, 6017, 0, 6, 22, 0, 6, 2, 0, 5961, 0\n0, 0, 0, 0, 0, 5742, 0, 0, 0,
0]\nnclass = d'),
Text(248.0, 543.5999999999999, 'O_3 <= 71.5\ngini = 0.482\nsamples = 788\n
value = [0, 0, 0, 39, 0, 0, 0, 0, 0, 0, 0, 406, 0, 0\n0, 0, 0, 0, 0, 788,
0, 0, 0, 0]\nnclass = t'),
Text(124.0, 181.19999999999982, 'gini = 0.453\nsamples = 441\nvalue = [0,
0, 0, 34, 0, 0, 0, 0, 0, 0, 0, 181, 0, 0\n0, 0, 0, 0, 0, 480, 0, 0, 0, 0]\n
```

Conclusion

Accuracy

```
In [134]: lr.score(x_train,y_train)
```

```
Out[134]: 0.3005822000263991
```

```
In [135]: rr.score(x_train,y_train)
```

```
Out[135]: 0.300578924013424
```

```
In [136]: la.score(x_train,y_train)
```

```
Out[136]: 0.04519637644952712
```

```
In [137]: en.score(x_test,y_test)
```

```
Out[137]: 0.15303243515292386
```

```
In [138]: logr.score(fs,target_vector)
```

```
Out[138]: 0.6612921669525443
```

```
In [139]: grid_search.best_score_
```

```
Out[139]: 0.6921165836260176
```

Linear Regression is suitable for this dataset